

Kathmandu Apache Spark Meetup

Jump Start with Apache® Spark™ 2.x on Databricks

<https://www.meetup.com/kathmandu-apache-spark>

<https://github.com/ganeshchand/kathmandu-apache-spark-meetup-introduction-to-spark-2.0>

\$ whoami

- Data & Analytics engineer @ Databricks
- Previously Developer @ Guidewire Software
- Past engineering roles at:
 - Accenture, SunGard, Hitachi Software
- ganesh@databricks.com
- <https://www.linkedin.com/in/chandganesh>
- @geechand

Agenda for the next 2.5 hours

Hour 1.5

- About KASM
- Get to know Databricks
- Overview of Spark
Fundamentals & Architecture
- What's New in Spark 2.0
- Unified APIs: SparkSessions,
SQL, DataFrames, Datasets...
- Workshop Notebook 1
- Break

Hour 1

- Introduction to DataFrames,
DataSets and Spark SQL
- Workshop Notebook 2
- Social hours

Kathmandu Apache Spark Meetup

A community of data practitioners and enthusiasts

Area of Interests:

- Functional programming (**scala**)
- Data Engineering (**Spark, Scala**)
- Data Science / Machine Learning (**Spark, Python, R**)
- Data Visualization (**D3.js**)
- Distributed Data Store (**Hadoop, S3, Redshift, Cassandra**)
- Cloud Platform (**Databricks, EMR**)
- Emerging tools & technologies (**Tenserflow, Apache Flink**)

<https://www.meetup.com/kathmandu-apache-spark>



Kathmandu Apache Spark Meetup

2015

- Introduction to Big Data and Apache Spark
- Live coding
- 55+ attendees
- Sponsored by RoyalePi
- Guest Speaker from France - Ludwine Probst



Get to know Databricks

- Get Databricks community edition <http://databricks.com/try-databricks>



WHY DATABRICKS PRODUCT APACHE SPARK SOLUTIONS CUSTOMERS TRAINING **TRY DATABRICKS**

Select a version to get started.

FULL-PLATFORM TRIAL

Run Apache Spark to work

- Unlimited clusters
- Notebooks, dashboards, production jobs, REST API
- Interactive guide to Spark and Databricks
- Deployed to your AWS VPC
- BI tool integration
- 14-day free trial (excludes AWS charges)

START TODAY

COMMUNITY EDITION

Learn Apache Spark

- Mini 5GB cluster
- Interactive notebooks and dashboards
- Public environment to share your work

START TODAY

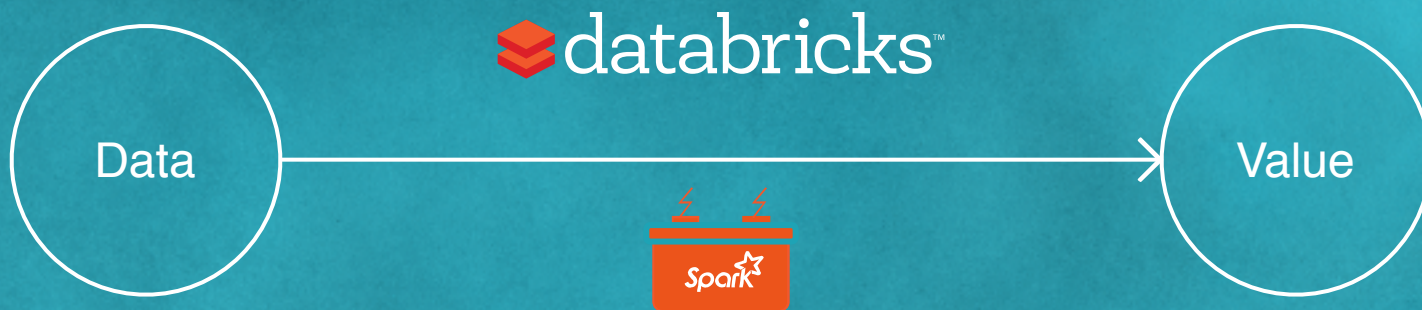
We are Databricks, the company behind Apache Spark



Founded by the creators
of Apache Spark in 2013

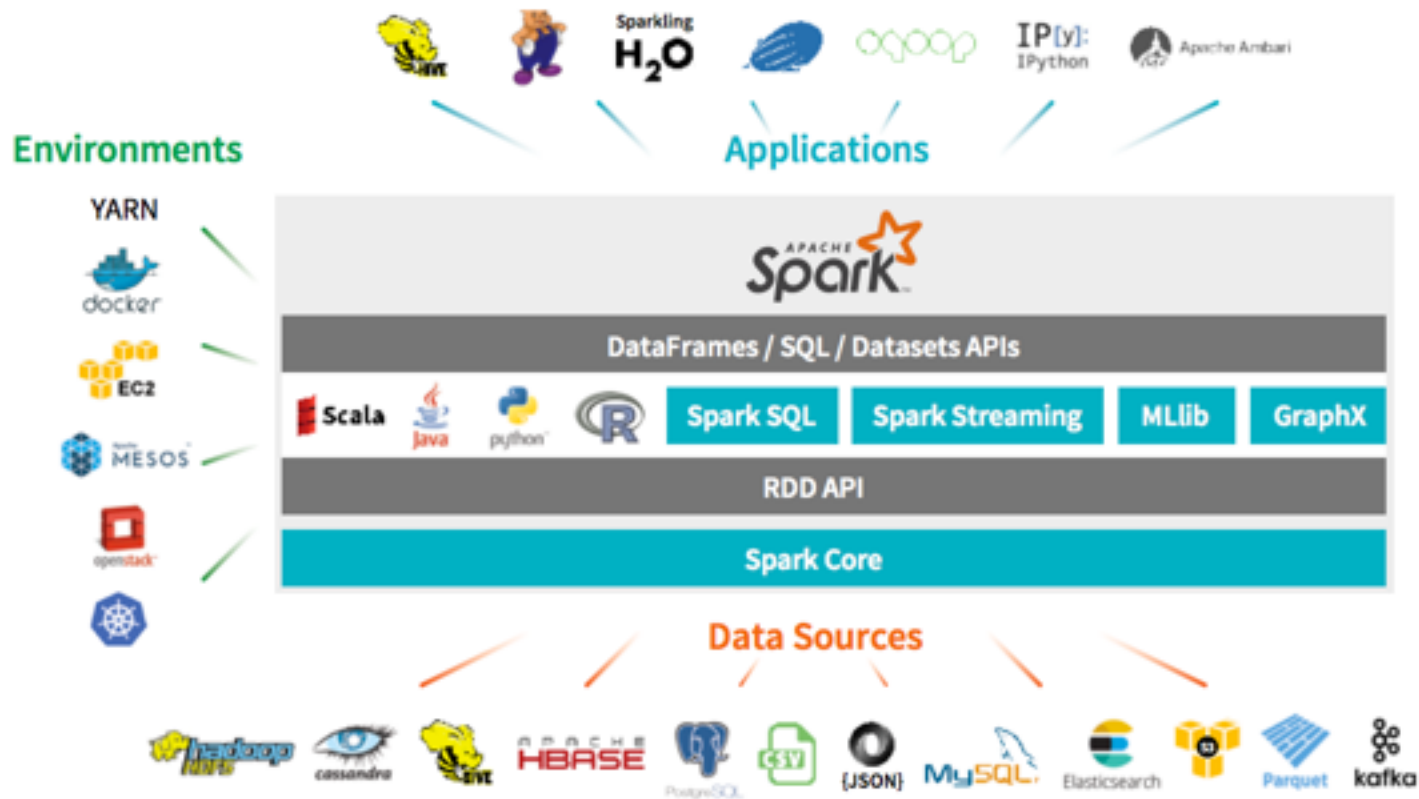
75%

Share of Spark code
contributed by
Databricks
in 2014



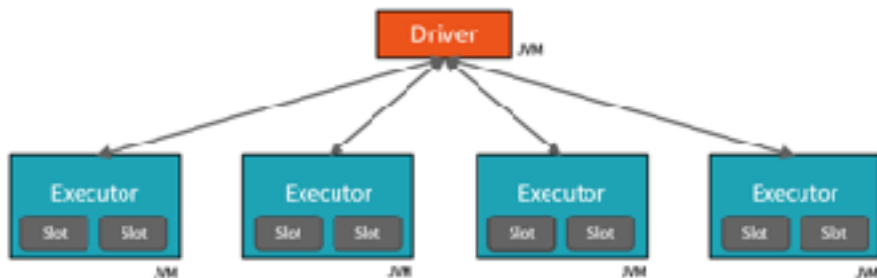
Created Databricks on top of Spark to make big data simple.

Unified engine across diverse workloads & environments



Apache Spark Architecture

Spark Physical Cluster

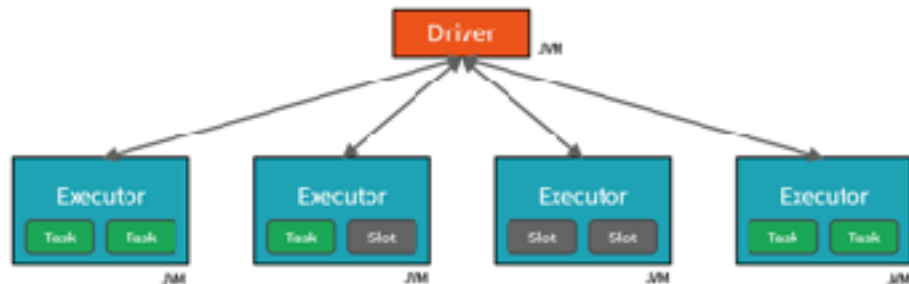


Deployments Modes

- Local
- Standalone
- YARN
- Mesos

Apache Spark Architecture

An Anatomy of an Application



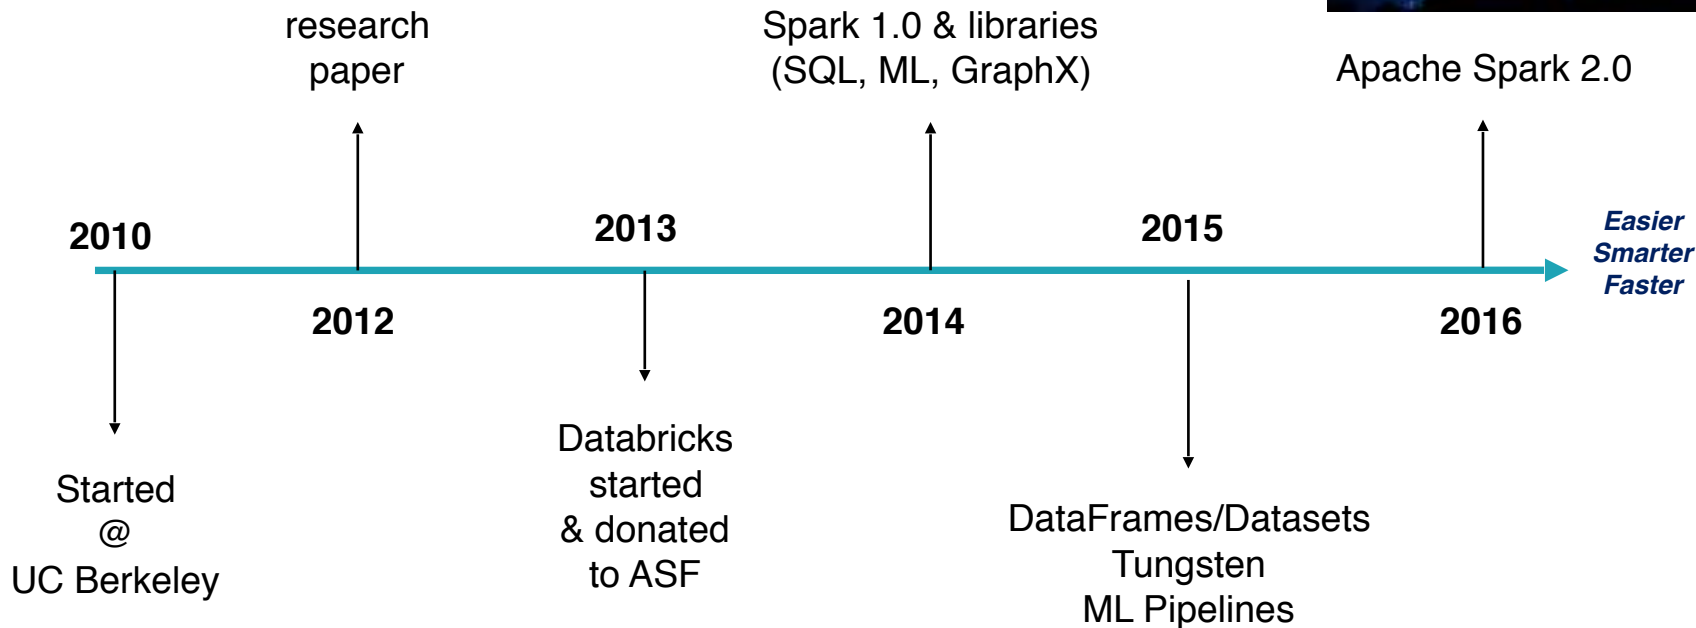
Spark Application

- Jobs
- Stages
- Tasks



How did we Get Here..?
Where we Going..?

A Brief History



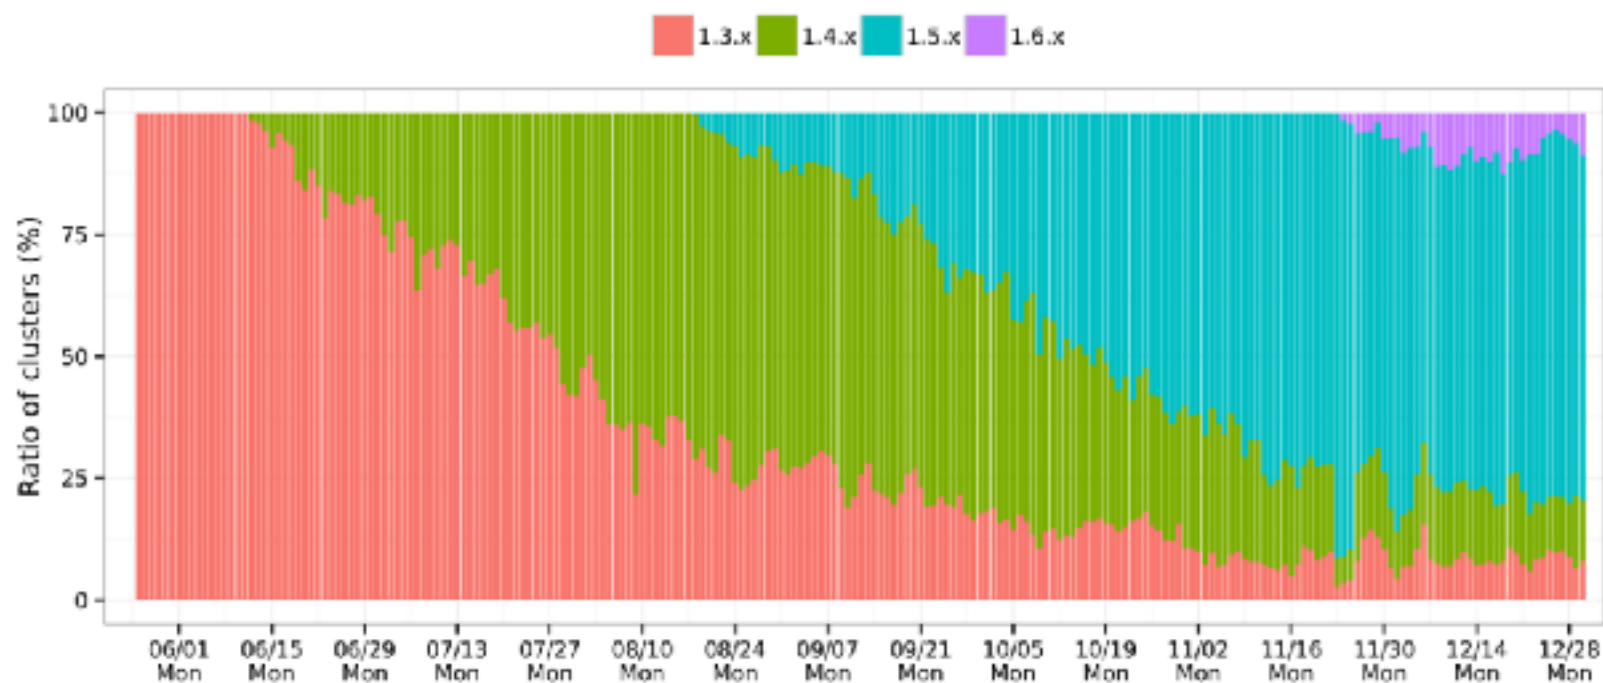
Apache Spark 2.0

- Steps to Bigger & Better Things....



Builds on all we learned in past 2 years

Percent of clusters by Spark version in 2015



Reynold Xin @rxin · Jan 14

The most interesting thing in Spark 2015 Year in Review post is the version dist over time: databricks.com/blog/2016/01/10

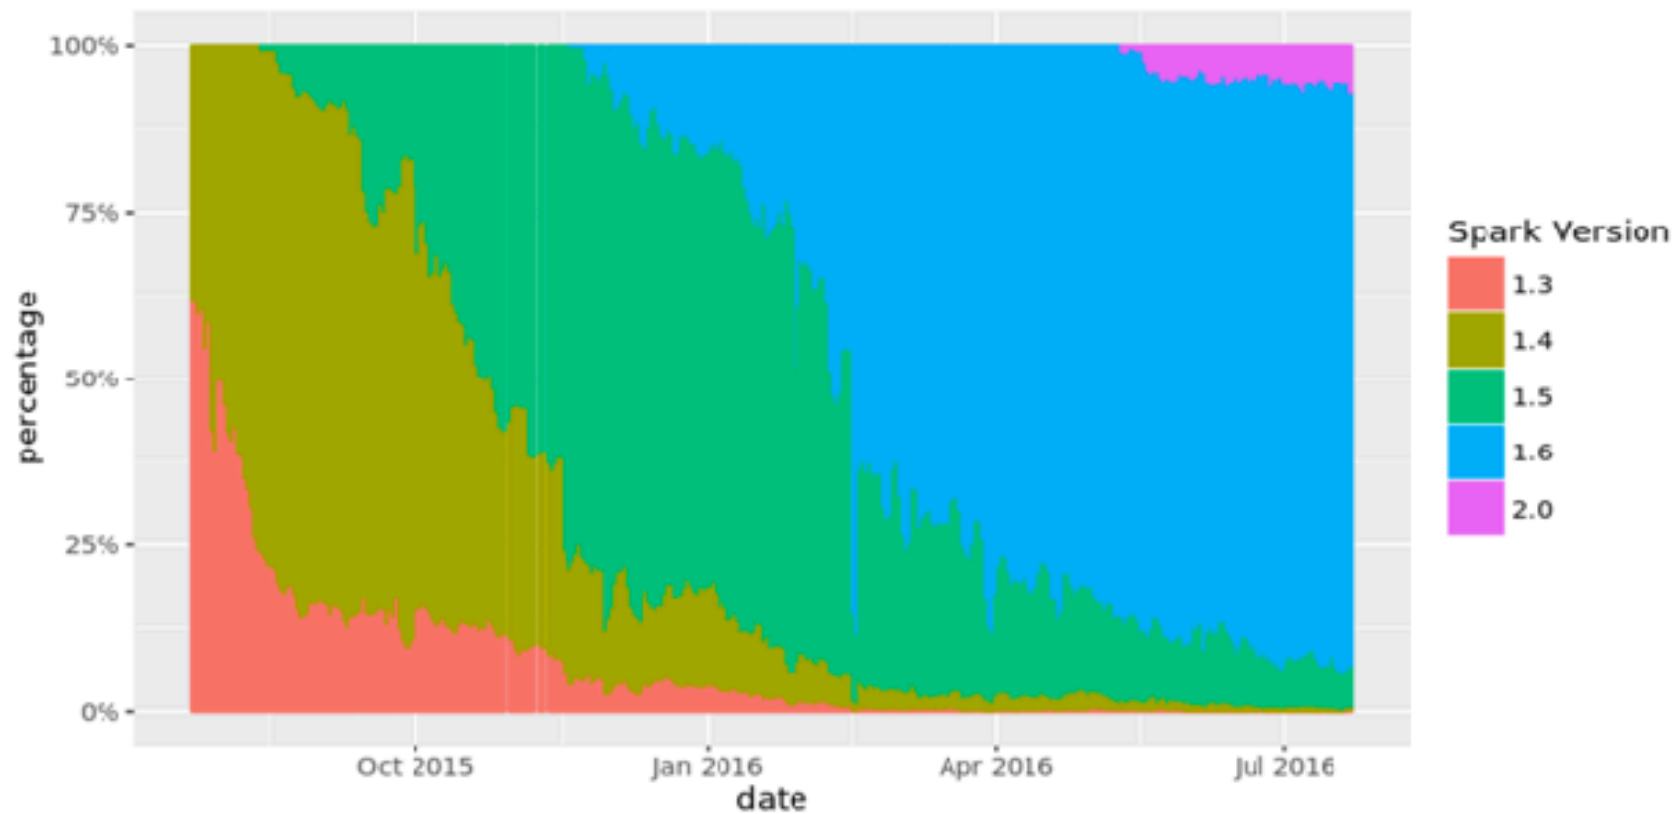


9



19





Apache Spark Usage over Time by Version

Major Features in Apache Spark 2.0



Unifying Datasets
and DataFrames &
SparkSessions

Easier



Tungsten Phase 2
speedups of 5-10x
& Catalyst Optimizer

Faster



Structured Streaming
real-time engine
on SQL / DataFrames

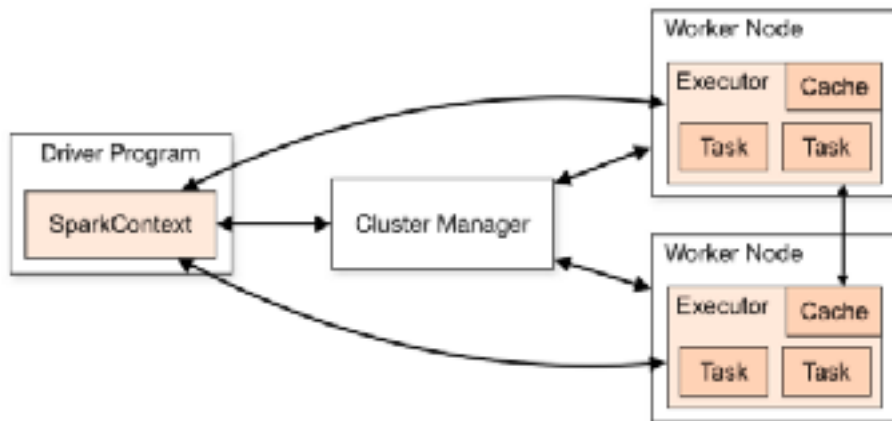
Smarter

Unified API Foundation for the
Future: Spark Sessions, Dataset,
DataFrame, MLlib, Structured
Streaming...

SparkSession – A Unified entry point to Spark

- SparkSession is the “SparkContext” for Dataset/DataFrame
 - Entry point for reading data and writing data
 - Working with metadata
 - Setting Spark Configuration
 - Driver uses for Cluster resource management

SparkSession vs SparkContext



SparkSessions Subsumes

- SparkContext
- SQLContext
- HiveContext
- StreamingContext
- SparkConf

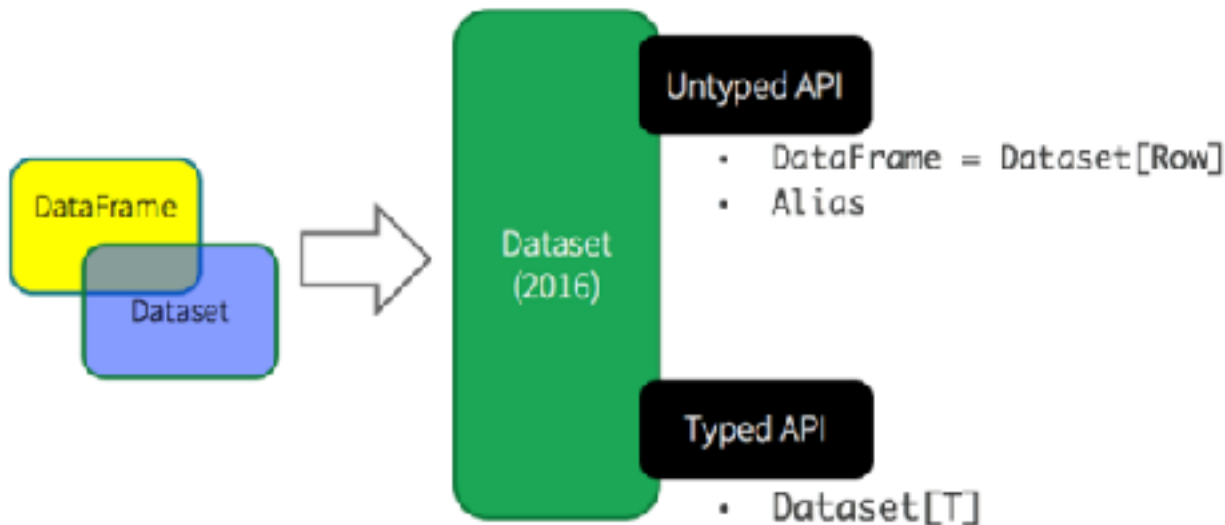
```
val warehouseLocation = "file:${system:user.dir}/spark-warehouse"

val spark = SparkSession
  .builder()
  .appName("SparkSessionZipsExample")
  .config("spark.sql.warehouse.dir", warehouseLocation)
  .enableHiveSupport()
  .getOrCreate()
```

Datasets and DataFrames



Unified Apache Spark 2.0 API



Long Term

- RDD will remain the low-level API in Spark
 - For control and certain type-safety in Java/Scala
- Datasets & DataFrames give richer semantics and optimizations
 - For semi-structured data and DSL like operations
 - New libraries will increasingly use these as interchange format
 - Examples: Structured Streaming, MLlib, GraphFrames
 - [A Tale of Three APIs: RDDs, DataFrames and Datasets](#)

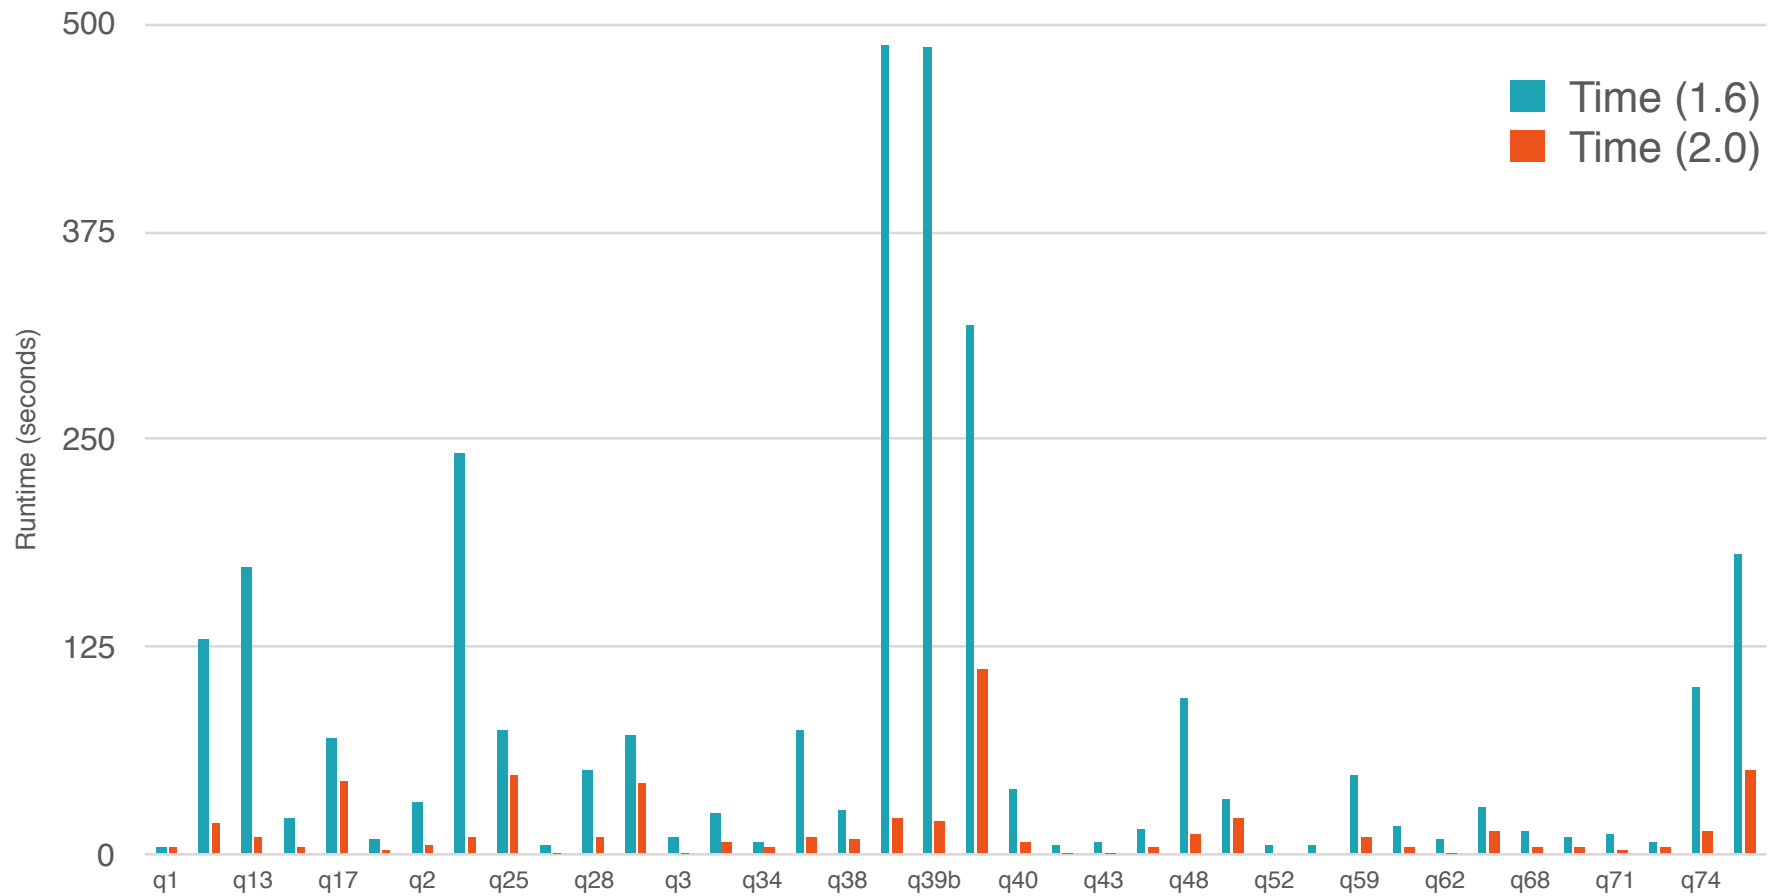
Other notable API improvements

- DataFrame-based ML pipeline API becoming the main MLlib API
- ML model & pipeline persistence with almost complete coverage
 - In all programming languages: Scala, Java, Python, R
- Improved R support
 - (Parallelizable) User-defined functions in R
 - Generalized Linear Models (GLMs), Naïve Bayes, Survival Regression, K-Means

Towards SQL 2003

- Today, Spark can run all 99 TPC-DS queries!
- New standard compliant parser (with good error messages!)
- Subqueries (correlated & uncorrelated)
- Approximate aggregate stats
 - <https://databricks.com/blog/2016/07/26/introducing-apache-spark-2-0.html>
 - <https://databricks.com/blog/2016/06/17/sql-subqueries-in-apache-spark-2-0.html>

Preliminary TPC-DS Spark 2.0 vs 1.6 – Lower is Better



How exactly do I learn and build Spark Apps?

- **Spark REPL** - ideal for learning and quick prototyping
- IDE Tool (**IntelliJ**) - ideal for library and end-to-end app development
- Interactive tools - ideal for exploring datasets and data storytelling.
 - **Databricks Notebook**
 - Zeppelin
 - Spark-Notebook
 - Jupyter-Scala

Workshop: Notebook on SparkSession

- Import Notebook into your Spark 2.0 Cluster
 - <http://dbricks.co/sswksh1>
 - <http://docs.databricks.com>
 - <http://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.sql.SparkSession>
- Familiarize your self with Databricks Notebook environment
- Work through each cell
 - CNTR + <return> / Shift + Return
- Try challenges
- Break...

DataFrames/Datasets & Spark SQL & Catalyst Optimizer

The not so secret truth...



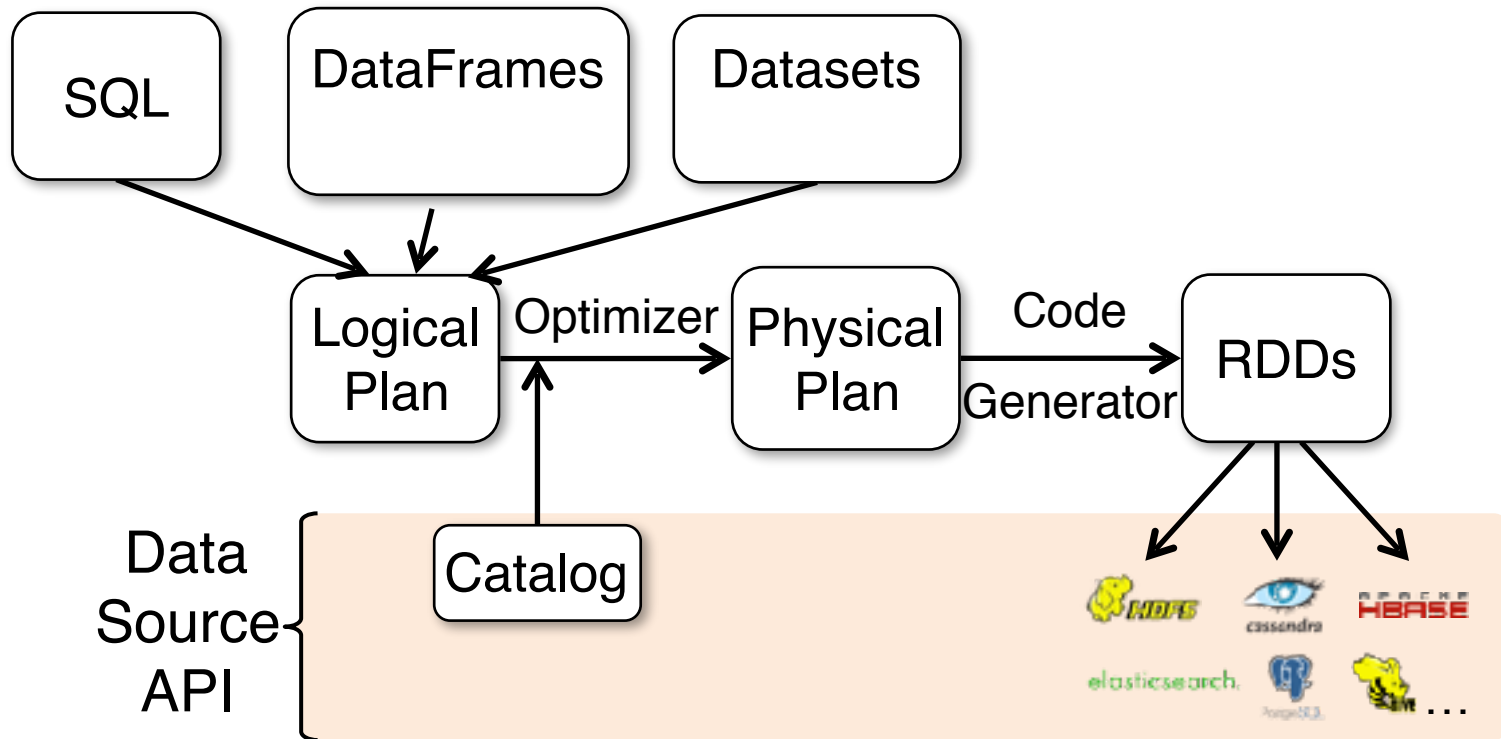
is not about SQL
is about more than
SQL

Spark SQL: The whole story

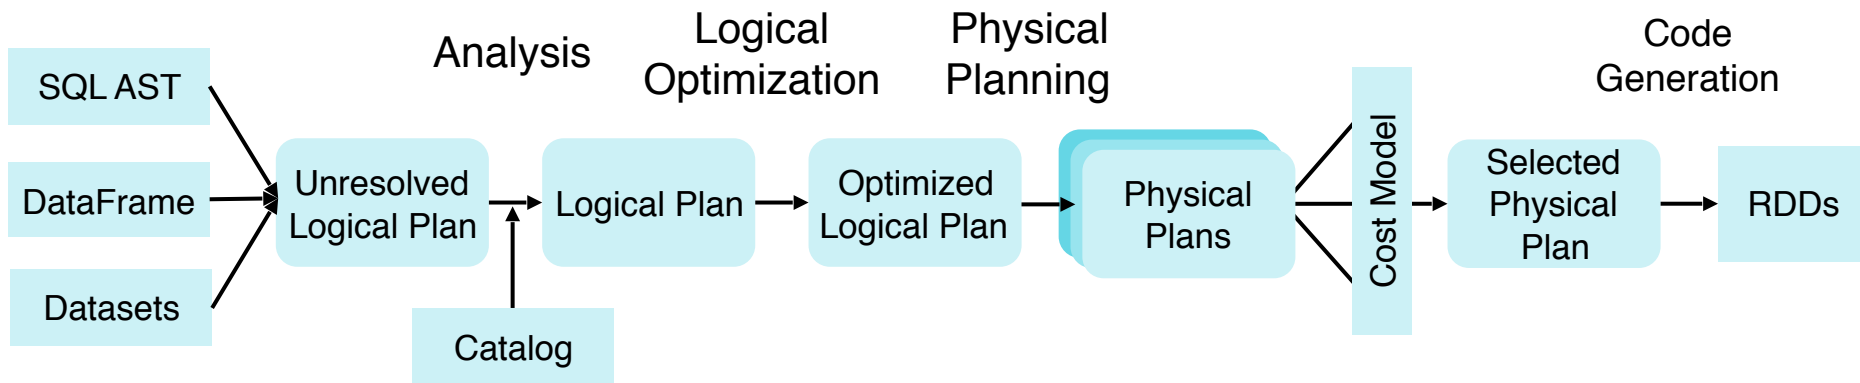
Is About Creating and Running Spark Programs Faster:

- Write less code
- Read less data
- Let the optimizer do the hard work

Spark SQL Architecture



Using Catalyst in Spark SQL



Analysis: analyzing a logical plan to resolve references

Logical Optimization: logical plan optimization

Physical Planning: Physical planning

Code Generation: Compile parts of the query to Java bytecode

Catalyst Optimizations

Logical Optimizations

- Push filter predicates down to data source, so irrelevant data can be skipped
- **Parquet:** skip entire blocks, turn comparisons on strings into cheaper integer comparisons via dictionary encoding
- **RDBMS:** reduce amount of data traffic by pushing predicates down

Create Physical Plan & generate JVM bytecode

- Catalyst compiles operations into physical plans for execution and generates JVM bytecode
- Intelligently choose between broadcast joins and shuffle joins to reduce network traffic
- **Lower level optimizations:** eliminate expensive object allocations and reduce virtual function calls


```
def add_demographics(events):
    u = sqlCtx.table("users")
    events \
        .join(u, events.user_id == u.user_id) \
        .withColumn("city", zipToCity(df.zip))
    events = add_demographics(sqlCtx.load("/data/events", "parquet"))
    training_data = events.where(events.city == "New York").select(events.timestamp).collect()
```

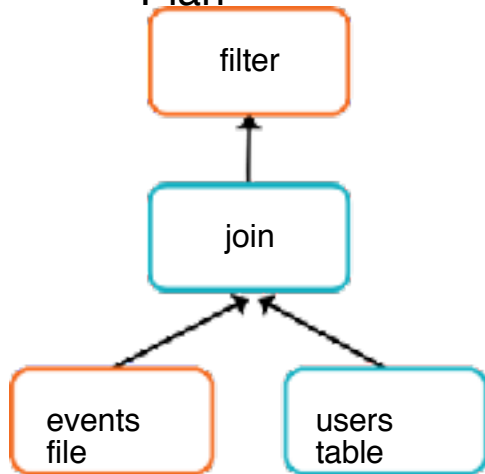
Load **partitioned** Hive table ←

Join on user_id

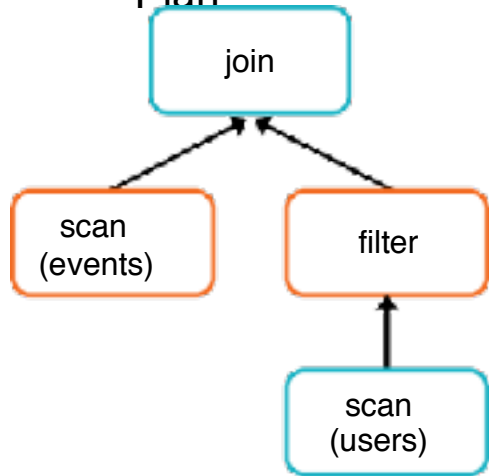
Run udf to add city column

←

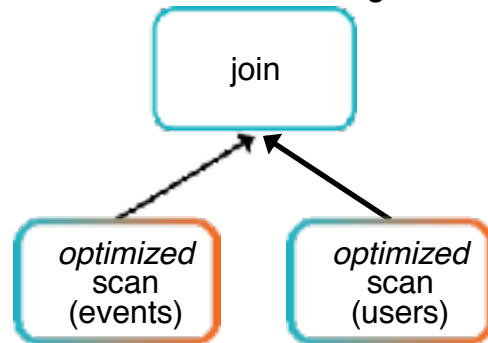
Logical Plan



Physical Plan



Physical Plan
with Predicate
Pushdown and
Column Pruning



Columns: Predicate pushdown

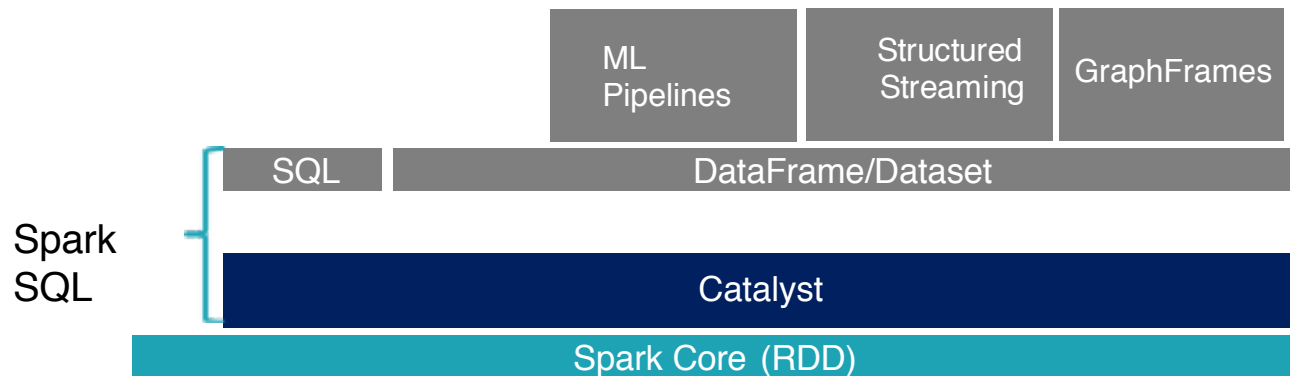
You Write

```
spark.read
  .format("jdbc")
  .option("url", "jdbc:postgresql:dbserver")
  .option("dbtable", "people")
  .load()
  .where($"name" === "michael")
```

Spark Translates
For Postgres

```
SELECT * FROM people WHERE name =  
'michael'
```

Foundational Spark 2.0 Components



Dataset Spark 2.0 APIs

Background: What is in an RDD?

- Dependencies
- Partitions (with optional locality info)
- Compute function: **Partition =>**
Iterator[T]

Opaque
Computation
& Opaque Data

Structured APIs In Spark



SQL

DataFrames

Datasets

Syntax
Errors

Runtime

Compile
Time

Compile
Time

Analysis
Errors

Runtime

Runtime

Compile
Time

Analysis errors are reported before a distributed job starts

Dataset API in Spark 2.0

- Typed interface over DataFrames / Tungsten

- `case class` Person(name: String, age: Int)
- `val` dataframe = spark.read.json("people.json")
- `val` ds: Dataset[Person] = dataframe.as[Person]
- ds.filter(p => p.name.startsWith("M"))
 .groupBy("name")
 .avg("age")

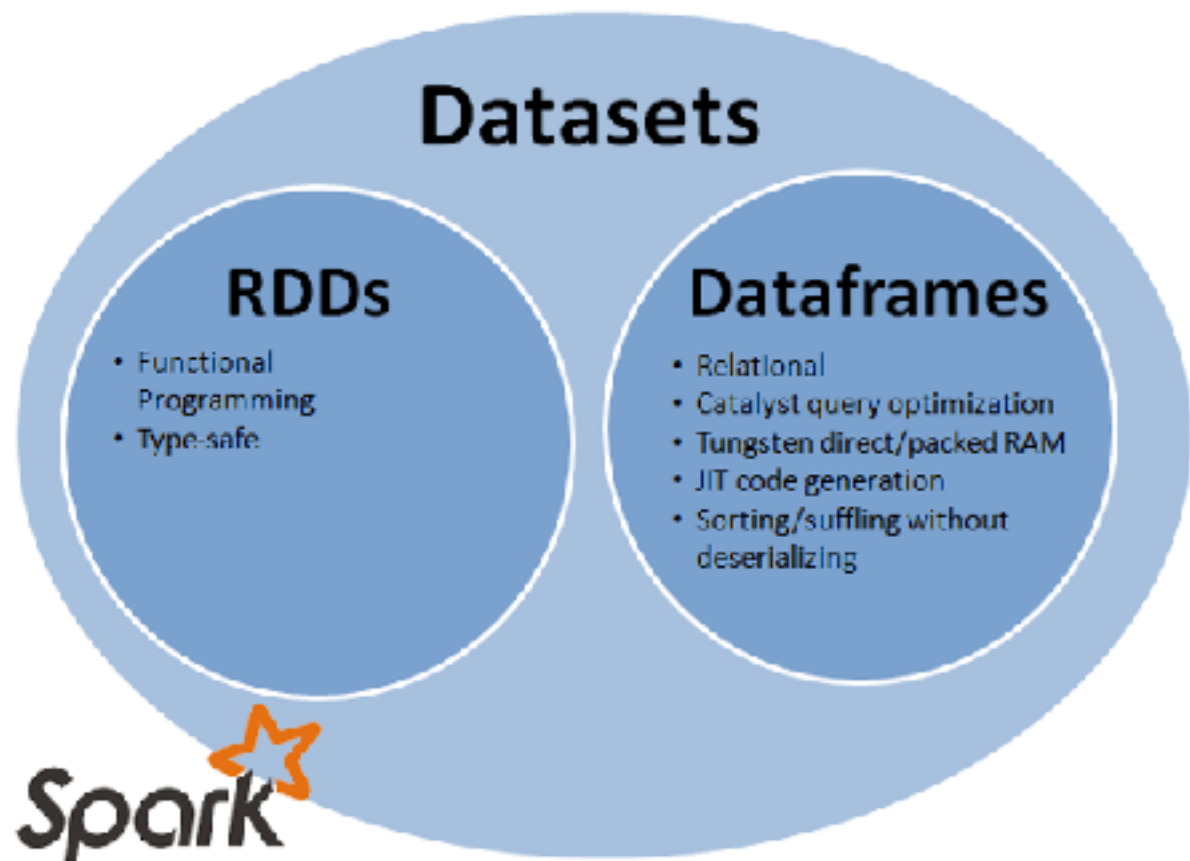
Datasets API

Type-safe:
operate on
domain objects
with compiled
lambda functions

```
val df = spark.read.json("people.json")

// Convert data to domain objects.
case class Person(name: String, age: Int)
val ds: Dataset[Person] = df.as[Person]
ds.filter(_.age > 30)

// Compute histogram of age by name.
val hist = ds.groupBy(_.name).mapGroups {
  case (name, people: Iter[Person]) =>
    val buckets = new Array[Int](10)
    people.map(_.age).foreach { a =>
      buckets(a / 10) += 1
    }
    (name, buckets)
}
```



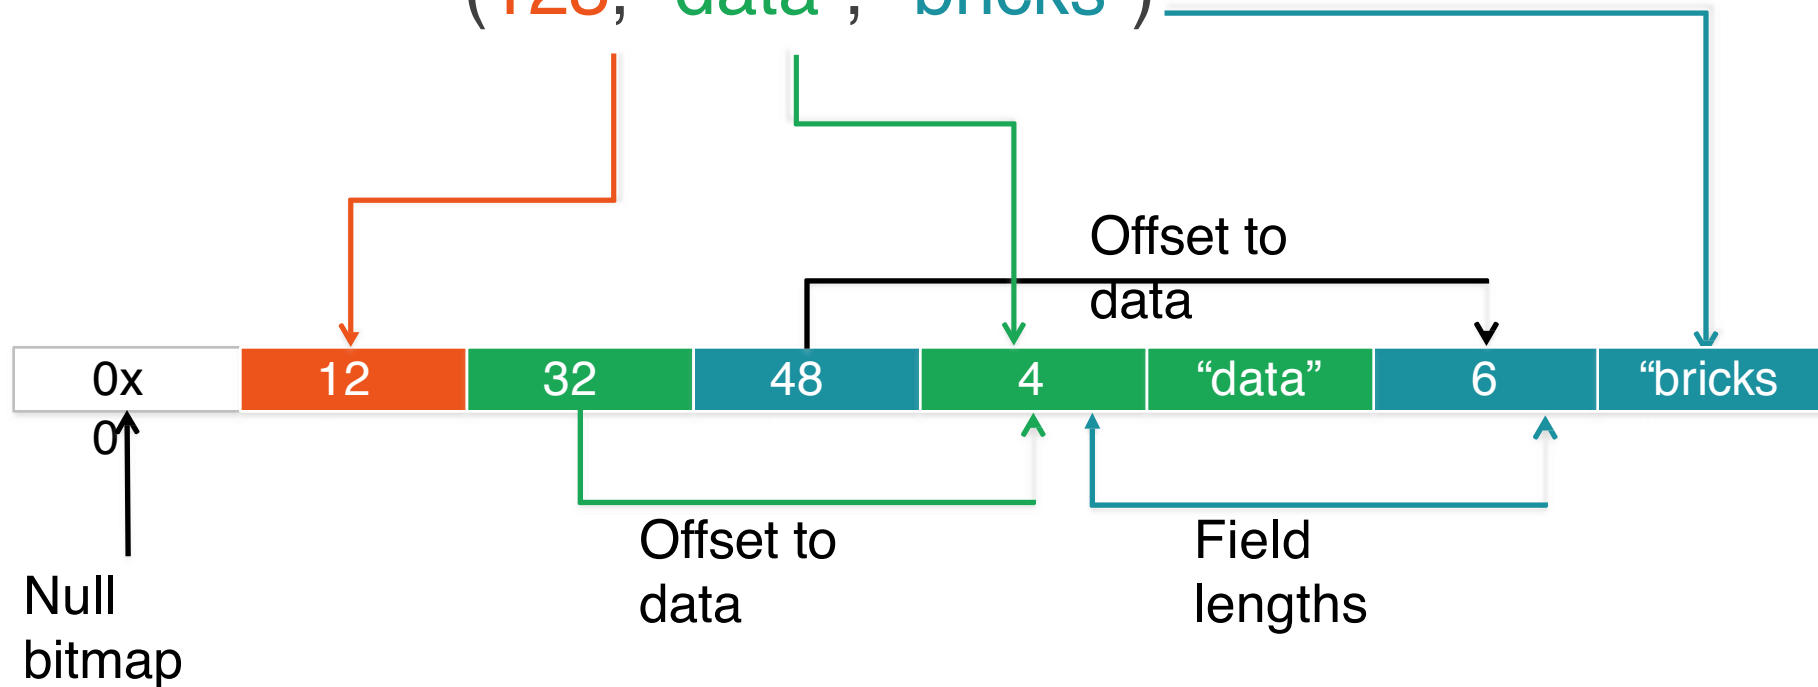
Project Tungsten II

Project Tungsten

- Substantially speed up execution by optimizing CPU efficiency, via: **SPARK-12795**
 - (1) Runtime code generation
 - (2) Exploiting cache locality
 - (3) Off-heap memory management

Tungsten's Compact Row Format

(123, "data", "bricks")



Encoders

Encoders translate between
domain objects and Spark's
internal format

JVM Object

MyClass(123, "data", "bricks")



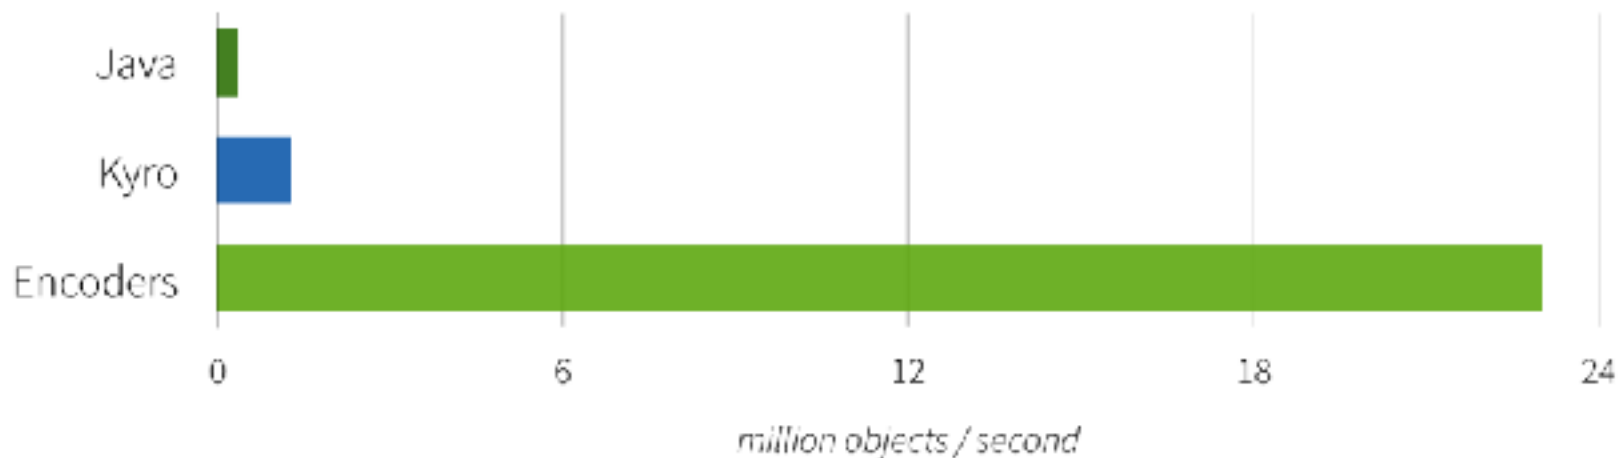
Internal Representation



0x0	123	32L	48L	4	"data"	6	"bricks"
-----	-----	-----	-----	---	--------	---	----------

Datasets: Lightning-fast Serialization with Encoders

Serialization / Deserialization Performance



Performance of Core Primitives

primitive	cost per row (single thread)	
	Spark 1.6	Spark 2.0
filter	15 ns	1.1 ns
sum w/o group	14 ns	0.9 ns
sum w/ group	79 ns	10.7 ns
hash join	115 ns	4.0 ns
sort (8 bit entropy)	620 ns	5.3 ns
sort (64 bit entropy)	620 ns	40 ns
sort-merge join	750 ns	700 ns

Project Tungsten: Bringing Apache Spark Closer to Bare Metal



by Reynold Xin and Joni Remigien
Posted in [Engineering](#) on April 28, 2015

In a previous [blog post](#), we looked back and surveyed performance improvements made to Apache Spark in the past year. In this post, we look forward and share with you the next chapter, which we are calling Project Tungsten. 2014 witnessed Spark setting the world record in large-scale sorting and saw major improvements across the entire engine from Python to SQL, to machine learning. Performance optimization, however, is a never-ending process.

Project Tungsten will be the largest change to Spark's execution engine since the project's inception. It focuses on substantially improving the efficiency of memory and CPU for Spark applications, to push performance closer to the limits of modern hardware. This effort includes three initiatives:

1. *Memory Management and Binary Processing*: leveraging application semantics to manage memory explicitly and eliminate the overhead of JVM object model and garbage collection
2. *Cache-aware computation*: algorithms and data structures to exploit memory hierarchy
3. *Code generation*: using code generation to exploit modern compilers and CPUs



SPARK SUMMIT 2016



Workshop: Notebook on DataFrames/ Datasets & Spark SQL

- Import Notebook into your Spark 2.0 Cluster
 - <http://dbricks.co/sswksh2A>
 - <http://dbricks.co/sswksh2>
 - <https://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.sql.Dataset>
- Work through each Notebook cell
- Try challenges
- Break..

Resources

- docs.databricks.com
- [Spark Programming Guide](#)
- [Structured Streaming Programming Guide](#)
- [Databricks Engineering Blogs](#)
- sparkhub.databricks.com
- <https://spark-packages.org/>

Do you have any questions
for my prepared answers?

